

Healthcare Project-02

March 26, 2022

1 Name: Sunil Pradhan

1.1 DOMAIN-HEALTHCARE

1.2 OBJECTIVE-

1.NIDDK (National Institute of Diabetes and Digestive and Kidney Diseases) research creates knowledge about and treatments for the most chronic, costly, and consequential diseases.

2.The dataset used in this project is originally from NIDDK. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset.

3.Build a model to accurately predict whether the patients in the dataset have diabetes or not.

1.3 Project Task: Week 1(Data Exploration)

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
[2]: df=pd.read_csv("health care diabetes.csv")
df.head()
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0

4 2.288 33 1

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                  768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                  768 non-null   int64
8   Outcome              768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[4]: df.describe()
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[5]: df.columns
```

```
[5]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
        'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
```

```
dtype='object')
```

```
[6]: df[df["Glucose"]==0]
```

```
[6]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
75             1         0             48             20         0  24.7
182            1         0             74             20        23  27.7
342            1         0             68             35         0  32.0
349            5         0             80             32         0  41.0
502            6         0             68             41         0  39.0


      DiabetesPedigreeFunction  Age  Outcome
75                        0.140   22         0
182                       0.299   21         0
342                       0.389   22         0
349                       0.346   37         1
502                       0.727   41         1
```

```
[7]: df["Glucose"].value_counts()
```

```
[7]: 99      17
100     17
111     14
129     14
125     14
..
191      1
177      1
44       1
62       1
190      1
Name: Glucose, Length: 136, dtype: int64
```


```
[8]: plt.figure(figsize=(10,7))
sns.histplot(df["Glucose"])
```

```
[8]: <AxesSubplot:xlabel='Glucose', ylabel='Count'>
```




```
[9]: plt.figure(figsize=(10,7))  
sns.histplot(df["BloodPressure"])
```

```
[9]: <AxesSubplot:xlabel='BloodPressure', ylabel='Count'>
```




```
[10]: plt.figure(figsize=(10,7))  
sns.histplot(df["BMI"])
```

```
[10]: <AxesSubplot:xlabel='BMI', ylabel='Count'>
```




```
[11]: plt.figure(figsize=(10,7))  
sns.histplot(df["SkinThickness"])
```

```
[11]: <AxesSubplot:xlabel='SkinThickness', ylabel='Count'>
```



```
[12]: plt.figure(figsize=(10,7))  
sns.histplot(df["Insulin"])
```

```
[12]: <AxesSubplot:xlabel='Insulin', ylabel='Count'>
```



```
[13]: df["Glucose"].mean()
```

```
[13]: 120.89453125
```

1.3.1 Repalce missing values i.e. 0

```
[14]: df["Glucose"] = df["Glucose"].replace(0, df["Glucose"].median())
df["BMI"] = df["BMI"].replace(0, df["BMI"].mean())
df["BloodPressure"] = df["BloodPressure"].replace(0, df["BloodPressure"].median())
df["Insulin"] = df["Insulin"].replace(0, df["Insulin"].median())
df["SkinThickness"] = df["SkinThickness"].replace(0, df["SkinThickness"].median())
```

```
[15]: df.head()
```

[15]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	30.5	33.6	
1	1	85	66	29	30.5	26.6	
2	8	183	64	23	30.5	23.3	
3	1	89	66	23	94.0	28.1	
4	0	137	40	35	168.0	43.1	

DiabetesPedigreeFunction Age Outcome

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[16]: df.tail()
```

```
[16]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
763           10     101           76           48     180.0  32.9
764           2     122           70           27      30.5  36.8
765           5     121           72           23     112.0  26.2
766           1     126           60           23      30.5  30.1
767           1      93           70           31      30.5  30.4

      DiabetesPedigreeFunction  Age  Outcome
763                0.171     63         0
764                0.340     27         0
765                0.245     30         0
766                0.349     47         1
767                0.315     23         0
```


1.3.2 There are integer and float data type variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables

```
[17]: df.dtypes.value_counts()
```

```
[17]: int64      6
float64      3
dtype: int64
```

```
[18]: plt.figure(figsize=(7,7))
sns.countplot(df.dtypes)
```


```
[18]: <AxesSubplot:ylabel='count'>
```



1.3.3 Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action

```
[19]: plt.figure(figsize=(10,7))  
sns.countplot(df["Outcome"])
```

```
[19]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```

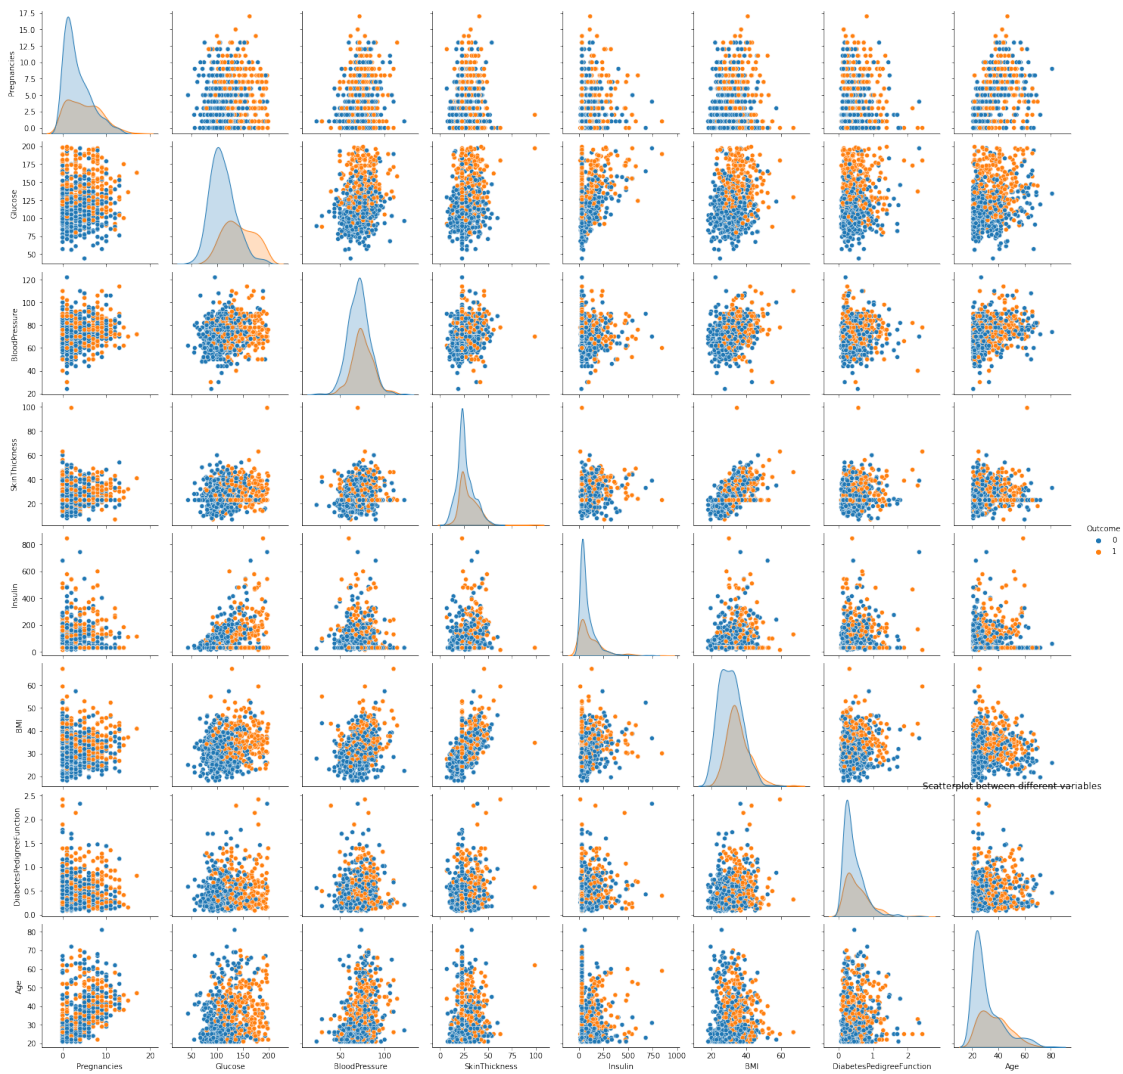


Since, the dataset is balanced. Thus, we don't apply random sampling technique for balancing the dataset.

1.4 Create scatter charts between the pair of variables to understand the relationships


```
[20]: sns.pairplot(data=df, hue="Outcome")  
plt.title("Scatterplot between different variables")
```

```
[20]: Text(0.5, 1.0, 'Scatterplot between different variables')
```



```
[21]: plt.figure(figsize=(10,7))
      sns.scatterplot("Glucose", "BloodPressure", data=df, hue="Outcome")
```

```
[21]: <AxesSubplot:xlabel='Glucose', ylabel='BloodPressure'>
```



1.5 Perform correlation analysis. Visually explore it using a heat map

```
[22]: df.corr()
```

```
[22]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.128213	0.208615	0.032568	
Glucose	0.128213	1.000000	0.218937	0.172143	
BloodPressure	0.208615	0.218937	1.000000	0.147809	
SkinThickness	0.032568	0.172143	0.147809	1.000000	
Insulin	-0.055697	0.357573	-0.028721	0.238188	
BMI	0.021546	0.231408	0.281129	0.546958	
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378	0.142977	
Age	0.544341	0.266909	0.324915	0.054514	
Outcome	0.221898	0.492782	0.165723	0.189065	


	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.055697	0.021546	-0.033523	
Glucose	0.357573	0.231408	0.137327	
BloodPressure	-0.028721	0.281129	-0.002378	
SkinThickness	0.238188	0.546958	0.142977	
Insulin	1.000000	0.189031	0.178029	
BMI	0.189031	1.000000	0.153508	

DiabetesPedigreeFunction	0.178029	0.153508	1.000000
Age	-0.015413	0.025748	0.033561
Outcome	0.148457	0.312254	0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.054514	0.189065
Insulin	-0.015413	0.148457
BMI	0.025748	0.312254
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[23]: plt.figure(figsize=(10,10))
      sns.heatmap(df.corr(),cmap="bwr",square=True,annot=True)
```

```
[23]: <AxesSubplot:>
```



1.6 Project Task: Week 2(Data modeling)

- 1.Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.
- 2.Apply an appropriate classification algorithm to build a model.
- 3.Compare various models with the results from KNN algorithm.
- 4.Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc.

```
[24]: df.head()
```

```
[24]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6     148             72             35     30.5  33.6
1           1      85             66             29     30.5  26.6
```

2	8	183	64	23	30.5	23.3
3	1	89	66	23	94.0	28.1
4	0	137	40	35	168.0	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[25]: x=df.drop("Outcome",axis=1)
      y=df[["Outcome"]]
```

```
[26]: x.head()
```

```
[26]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0             6      148            72           35      30.5  33.6
1             1       85            66           29      30.5  26.6
2             8      183            64           23      30.5  23.3
3             1       89            66           23      94.0  28.1
4             0      137            40           35     168.0  43.1
```

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

```
[27]: y.head()
```

```
[27]: Outcome
0      1
1      0
2      1
3      0
4      1
```

```
[28]: from sklearn.model_selection import train_test_split
```

```
[29]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪25,random_state=25)
```

```
[30]: x_train.shape,x_test.shape
```

```
[30]: ((576, 8), (192, 8))
```



```
[31]: y_train.shape,y_test.shape
```

```
[31]: ((576, 1), (192, 1))
```

```
[32]: x_train.head()
```

```
[32]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
427             1     181             64             30    180.0  34.1
133             8      84             74             31     30.5  38.3
262             4      95             70             32     30.5  32.1
358            12      88             74             40     54.0  35.3
489             8     194             80             23     30.5  26.1

      DiabetesPedigreeFunction  Age
427                   0.328    38
133                   0.457    39
262                   0.612    24
358                   0.378    48
489                   0.551    67
```

```
[33]: from sklearn.preprocessing import StandardScaler
```

```
[34]: scaler=StandardScaler()
x_train=scaler.fit_transform(x_train)
x_test=scaler.transform(x_test)
```

```
[35]: x_train
```

```
[35]: array([[ -0.82827947,  1.9049318 , -0.72994462, ...,  0.2534439 ,
        -0.42508999,  0.37784439],
        [ 1.24216246, -1.24181301,  0.10494041, ...,  0.86654809,
        -0.0415753 ,  0.46339407],
        [ 0.05905278, -0.88496566, -0.2290136 , ..., -0.03851047,
        0.41923693, -0.81985104],
        ...,
        [ 0.05905278, -0.26859296, -0.72994462, ..., -0.50563747,
        -1.02563238, -0.81985104],
        [-0.23672464, -0.23615229, -0.56296761, ...,  0.83735265,
        -0.95428081, -0.47765234],
        [-0.23672464,  1.54808445, -0.72994462, ...,  0.31183478,
        -0.34184649, -0.306553  ]])
```

1.7 Logistic regression

```
[36]: from sklearn.linear_model import LogisticRegression
```

```
[37]: log_model=LogisticRegression()  
log_model.fit(x_train,y_train)
```

```
[37]: LogisticRegression()
```

```
[38]: log_pred=log_model.predict(x_test)  
log_pred
```

```
[38]: array([1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,  
        0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0,  
        0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,  
        0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1,  
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1,  
        1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,  
        0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0,  
        0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
        0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1], dtype=int64)
```

```
[39]: from sklearn.metrics import   
      ↪ accuracy_score, confusion_matrix, classification_report, roc_curve, auc, RocCurveDisplay
```

```
[40]: accuracy_score(log_pred,y_test)
```

```
[40]: 0.8020833333333334
```

```
[41]: accuracy_score(log_model.predict(x_train),y_train)
```

```
[41]: 0.765625
```

```
[42]: confusion_matrix(log_pred,y_test)
```

```
[42]: array([[115,  23],  
        [ 15,  39]], dtype=int64)
```

1.8 Sensitivity=TP/(TP+FN)

1.9 Specificity=TN/(TN+FP)

```
[43]: print("Sensitivity is ",115/(115+39))  
      print("Specificity is ",23/(23+15))
```

Sensitivity is 0.7467532467532467

Specificity is 0.6052631578947368

```
[44]: print(classification_report(log_pred,y_test))
```

```
precision    recall  f1-score   support
```

0	0.88	0.83	0.86	138
1	0.63	0.72	0.67	54
accuracy				0.80
macro avg	0.76	0.78	0.77	192
weighted avg	0.81	0.80	0.81	192

```
[45]: log_model.predict_proba(x_test)
```

```
[45]: array([[0.49718064, 0.50281936],
 [0.5184736 , 0.4815264 ],
 [0.63650222, 0.36349778],
 [0.23770204, 0.76229796],
 [0.68861673, 0.31138327],
 [0.71208456, 0.28791544],
 [0.9387175 , 0.0612825 ],
 [0.79443084, 0.20556916],
 [0.92992385, 0.07007615],
 [0.90652177, 0.09347823],
 [0.33810021, 0.66189979],
 [0.85035462, 0.14964538],
 [0.70225282, 0.29774718],
 [0.48135433, 0.51864567],
 [0.77123986, 0.22876014],
 [0.30883256, 0.69116744],
 [0.81033777, 0.18966223],
 [0.87974871, 0.12025129],
 [0.71237312, 0.28762688],
 [0.64732719, 0.35267281],
 [0.80837603, 0.19162397],
 [0.89390753, 0.10609247],
 [0.92169482, 0.07830518],
 [0.57197108, 0.42802892],
 [0.9186062 , 0.0813938 ],
 [0.66550181, 0.33449819],
 [0.67938683, 0.32061317],
 [0.238574 , 0.761426 ],
 [0.78020378, 0.21979622],
 [0.83503884, 0.16496116],
 [0.66585115, 0.33414885],
 [0.35056697, 0.64943303],
 [0.87144477, 0.12855523],
 [0.12191586, 0.87808414],
 [0.4654908 , 0.5345092 ],
 [0.75957272, 0.24042728],
 [0.11965296, 0.88034704],
```

[0.60990487, 0.39009513],
[0.68606143, 0.31393857],
[0.89322221, 0.10677779],
[0.51946661, 0.48053339],
[0.92167769, 0.07832231],
[0.33351575, 0.66648425],
[0.66547305, 0.33452695],
[0.98254506, 0.01745494],
[0.8455861 , 0.1544139],
[0.77630737, 0.22369263],
[0.91105057, 0.08894943],
[0.84328445, 0.15671555],
[0.94660385, 0.05339615],
[0.23362354, 0.76637646],
[0.20330587, 0.79669413],
[0.333783 , 0.666217],
[0.61136828, 0.38863172],
[0.23513457, 0.76486543],
[0.82400676, 0.17599324],
[0.16884429, 0.83115571],
[0.93137853, 0.06862147],
[0.53440194, 0.46559806],
[0.68716648, 0.31283352],
[0.52076909, 0.47923091],
[0.73632975, 0.26367025],
[0.89867379, 0.10132621],
[0.70994113, 0.29005887],
[0.85601364, 0.14398636],
[0.09284684, 0.90715316],
[0.96629337, 0.03370663],
[0.3066029 , 0.6933971],
[0.86865154, 0.13134846],
[0.97029482, 0.02970518],
[0.8342464 , 0.1657536],
[0.4324549 , 0.5675451],
[0.80664318, 0.19335682],
[0.74969416, 0.25030584],
[0.18987813, 0.81012187],
[0.66750145, 0.33249855],
[0.11518964, 0.88481036],
[0.12647019, 0.87352981],
[0.88018424, 0.11981576],
[0.85975988, 0.14024012],
[0.906061 , 0.093939],
[0.92246813, 0.07753187],
[0.47639989, 0.52360011],
[0.1427741 , 0.8572259],

[0.76896699, 0.23103301],
[0.83572397, 0.16427603],
[0.06927278, 0.93072722],
[0.33457114, 0.66542886],
[0.93759669, 0.06240331],
[0.84767778, 0.15232222],
[0.76947422, 0.23052578],
[0.55218153, 0.44781847],
[0.96335471, 0.03664529],
[0.88596527, 0.11403473],
[0.85699598, 0.14300402],
[0.77839507, 0.22160493],
[0.86506761, 0.13493239],
[0.86802452, 0.13197548],
[0.93295922, 0.06704078],
[0.77638548, 0.22361452],
[0.37226787, 0.62773213],
[0.73133942, 0.26866058],
[0.77146353, 0.22853647],
[0.23802194, 0.76197806],
[0.87201793, 0.12798207],
[0.83925985, 0.16074015],
[0.70895277, 0.29104723],
[0.55420225, 0.44579775],
[0.26575849, 0.73424151],
[0.30988722, 0.69011278],
[0.34069358, 0.65930642],
[0.36265149, 0.63734851],
[0.93383866, 0.06616134],
[0.78830765, 0.21169235],
[0.95814728, 0.04185272],
[0.49761356, 0.50238644],
[0.72859611, 0.27140389],
[0.12896551, 0.87103449],
[0.69571374, 0.30428626],
[0.86062231, 0.13937769],
[0.95517846, 0.04482154],
[0.72404898, 0.27595102],
[0.8363083, 0.1636917],
[0.41996575, 0.58003425],
[0.66792736, 0.33207264],
[0.83436533, 0.16563467],
[0.42751497, 0.57248503],
[0.73467252, 0.26532748],
[0.93292395, 0.06707605],
[0.45311911, 0.54688089],
[0.90916868, 0.09083132],

[0.62650398, 0.37349602],
[0.75332714, 0.24667286],
[0.98285703, 0.01714297],
[0.80157566, 0.19842434],
[0.74187644, 0.25812356],
[0.68724046, 0.31275954],
[0.7099196 , 0.2900804],
[0.72973092, 0.27026908],
[0.83523309, 0.16476691],
[0.36691741, 0.63308259],
[0.57148589, 0.42851411],
[0.28444974, 0.71555026],
[0.95722719, 0.04277281],
[0.96079177, 0.03920823],
[0.14725599, 0.85274401],
[0.73986931, 0.26013069],
[0.90349332, 0.09650668],
[0.77986338, 0.22013662],
[0.74419975, 0.25580025],
[0.26370149, 0.73629851],
[0.53548182, 0.46451818],
[0.25734521, 0.74265479],
[0.91548717, 0.08451283],
[0.88917167, 0.11082833],
[0.69932255, 0.30067745],
[0.60597638, 0.39402362],
[0.78584319, 0.21415681],
[0.21870737, 0.78129263],
[0.09375074, 0.90624926],
[0.58195365, 0.41804635],
[0.9475746 , 0.0524254],
[0.09090876, 0.90909124],
[0.80440245, 0.19559755],
[0.95809381, 0.04190619],
[0.2281093 , 0.7718907],
[0.86876951, 0.13123049],
[0.96019502, 0.03980498],
[0.89661361, 0.10338639],
[0.98674765, 0.01325235],
[0.64401438, 0.35598562],
[0.62612338, 0.37387662],
[0.18944878, 0.81055122],
[0.91492112, 0.08507888],
[0.91848601, 0.08151399],
[0.77267372, 0.22732628],
[0.72794938, 0.27205062],
[0.48860969, 0.51139031],

```
[0.6171518 , 0.3828482 ],
[0.44535515, 0.55464485],
[0.23443856, 0.76556144],
[0.74161218, 0.25838782],
[0.8793028 , 0.1206972 ],
[0.89664476, 0.10335524],
[0.18860671, 0.81139329],
[0.70937628, 0.29062372],
[0.92169631, 0.07830369],
[0.33210235, 0.66789765],
[0.10412118, 0.89587882],
[0.93956879, 0.06043121],
[0.5617854 , 0.4382146 ],
[0.04899981, 0.95100019]])
```

```
[46]: log_pred_proba=log_model.predict_proba(x_test)[: ,1]
      log_pred_proba
```

```
[46]: array([0.50281936, 0.4815264 , 0.36349778, 0.76229796, 0.31138327,
0.28791544, 0.0612825 , 0.20556916, 0.07007615, 0.09347823,
0.66189979, 0.14964538, 0.29774718, 0.51864567, 0.22876014,
0.69116744, 0.18966223, 0.12025129, 0.28762688, 0.35267281,
0.19162397, 0.10609247, 0.07830518, 0.42802892, 0.0813938 ,
0.33449819, 0.32061317, 0.761426 , 0.21979622, 0.16496116,
0.33414885, 0.64943303, 0.12855523, 0.87808414, 0.5345092 ,
0.24042728, 0.88034704, 0.39009513, 0.31393857, 0.10677779,
0.48053339, 0.07832231, 0.66648425, 0.33452695, 0.01745494,
0.1544139 , 0.22369263, 0.08894943, 0.15671555, 0.05339615,
0.76637646, 0.79669413, 0.666217 , 0.38863172, 0.76486543,
0.17599324, 0.83115571, 0.06862147, 0.46559806, 0.31283352,
0.47923091, 0.26367025, 0.10132621, 0.29005887, 0.14398636,
0.90715316, 0.03370663, 0.6933971 , 0.13134846, 0.02970518,
0.1657536 , 0.5675451 , 0.19335682, 0.25030584, 0.81012187,
0.33249855, 0.88481036, 0.87352981, 0.11981576, 0.14024012,
0.093939 , 0.07753187, 0.52360011, 0.8572259 , 0.23103301,
0.16427603, 0.93072722, 0.66542886, 0.06240331, 0.15232222,
0.23052578, 0.44781847, 0.03664529, 0.11403473, 0.14300402,
0.22160493, 0.13493239, 0.13197548, 0.06704078, 0.22361452,
0.62773213, 0.26866058, 0.22853647, 0.76197806, 0.12798207,
0.16074015, 0.29104723, 0.44579775, 0.73424151, 0.69011278,
0.65930642, 0.63734851, 0.06616134, 0.21169235, 0.04185272,
0.50238644, 0.27140389, 0.87103449, 0.30428626, 0.13937769,
0.04482154, 0.27595102, 0.1636917 , 0.58003425, 0.33207264,
0.16563467, 0.57248503, 0.26532748, 0.06707605, 0.54688089,
0.09083132, 0.37349602, 0.24667286, 0.01714297, 0.19842434,
0.25812356, 0.31275954, 0.2900804 , 0.27026908, 0.16476691,
0.63308259, 0.42851411, 0.71555026, 0.04277281, 0.03920823,
```

```

0.85274401, 0.26013069, 0.09650668, 0.22013662, 0.25580025,
0.73629851, 0.46451818, 0.74265479, 0.08451283, 0.11082833,
0.30067745, 0.39402362, 0.21415681, 0.78129263, 0.90624926,
0.41804635, 0.0524254 , 0.90909124, 0.19559755, 0.04190619,
0.7718907 , 0.13123049, 0.03980498, 0.10338639, 0.01325235,
0.35598562, 0.37387662, 0.81055122, 0.08507888, 0.08151399,
0.22732628, 0.27205062, 0.51139031, 0.3828482 , 0.55464485,
0.76556144, 0.25838782, 0.1206972 , 0.10335524, 0.81139329,
0.29062372, 0.07830369, 0.66789765, 0.89587882, 0.06043121,
0.4382146 , 0.95100019])

```

```

[47]: fpr,tpr,thresholds=roc_curve(y_test,log_pred_proba)
print(fpr)
print("\n")
print(tpr)
print("\n")
print(thresholds)

```

```

[0.          0.          0.00769231 0.00769231 0.03076923 0.03076923
 0.03846154 0.03846154 0.05384615 0.05384615 0.06153846 0.06153846
 0.08461538 0.08461538 0.1          0.1          0.10769231 0.10769231
 0.11538462 0.11538462 0.16923077 0.16923077 0.17692308 0.17692308
 0.20769231 0.20769231 0.23076923 0.23076923 0.25384615 0.25384615
 0.28461538 0.28461538 0.29230769 0.29230769 0.36923077 0.36923077
 0.56923077 0.56923077 0.57692308 0.57692308 0.6          0.6
 0.62307692 0.62307692 0.63846154 0.63846154 0.66153846 0.66153846
 0.88461538 0.88461538 1.          ]

```

```

[0.          0.01612903 0.01612903 0.20967742 0.20967742 0.32258065
 0.32258065 0.43548387 0.43548387 0.48387097 0.48387097 0.51612903
 0.51612903 0.56451613 0.56451613 0.58064516 0.58064516 0.61290323
 0.61290323 0.66129032 0.66129032 0.69354839 0.69354839 0.72580645
 0.72580645 0.74193548 0.74193548 0.79032258 0.79032258 0.80645161
 0.80645161 0.82258065 0.82258065 0.85483871 0.85483871 0.87096774
 0.87096774 0.88709677 0.88709677 0.91935484 0.91935484 0.93548387
 0.93548387 0.9516129 0.9516129 0.96774194 0.96774194 0.98387097
 0.98387097 1.          1.          ]

```

```

[1.95100019 0.95100019 0.93072722 0.83115571 0.81012187 0.76229796
 0.76197806 0.69116744 0.66789765 0.66542886 0.66189979 0.64943303
 0.62773213 0.5675451 0.54688089 0.5345092 0.52360011 0.51139031
 0.50281936 0.48053339 0.42851411 0.41804635 0.39402362 0.38863172
 0.36349778 0.35598562 0.33449819 0.33207264 0.31283352 0.31275954
 0.29774718 0.29104723 0.29062372 0.29005887 0.26013069 0.25838782
 0.16563467 0.16496116 0.16476691 0.1636917 0.1544139 0.15232222

```




```
0.14300402 0.14024012 0.13493239 0.13197548 0.12855523 0.12798207
0.06240331 0.0612825 0.01325235]
```

```
[48]: roc_auc=auc(fpr,tpr)
      roc_auc
```

```
[48]: 0.8388337468982631
```

```
[49]: display=RocCurveDisplay(fpr=fpr,tpr=tpr,roc_auc=roc_auc,estimator_name="example_
      ↪estimator")
      display.plot()
      plt.plot(fpr,fpr,"r--",color="green")
```

```
[49]: [<matplotlib.lines.Line2D at 0x1ba8fa7bbe0>]
```



1.10 Support vector machine

```
[50]: from sklearn.svm import SVC
```

```
[51]: svc_model=SVC(probability=True)
      svc_model.fit(x_train,y_train)
      svc_pred=svc_model.predict(x_test)
      svc_pred
```

```
[51]: array([0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1,
          0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0,
          1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0,
          0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
          0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1], dtype=int64)
```

```
[52]: accuracy_score(svc_pred,y_test)
```

```
[52]: 0.78125
```

```
[53]: confusion_matrix(svc_pred,y_test)
```

```
[53]: array([[116, 28],
          [ 14, 34]], dtype=int64)
```

```
[54]: print("Sensitivity is ",116/(116+34))
```

```
Sensitivity is 0.7733333333333333
```

```
[55]: print("Specificity is ",25/(28+14))
```


```
Specificity is 0.5952380952380952
```

```
[56]: print(classification_report(svc_pred,y_test))
```

	precision	recall	f1-score	support
0	0.89	0.81	0.85	144
1	0.55	0.71	0.62	48
accuracy			0.78	192
macro avg	0.72	0.76	0.73	192
weighted avg	0.81	0.78	0.79	192

```
[57]: svc_pred_proba=svc_model.predict_proba(x_test)[:,-1]
fpr,tpr,thresholds=roc_curve(y_test,svc_pred_proba)
roc_auc=auc(fpr,tpr)
display=RocCurveDisplay(fpr=fpr,tpr=tpr,roc_auc=roc_auc,estimator_name="example_
→estimator")
display.plot()
plt.plot(fpr,fpr,"r--",color="red")
```

```
[57]: [<matplotlib.lines.Line2D at 0x1ba8fad6a90>]
```



1.11 Decision tree

```
[58]: from sklearn.tree import DecisionTreeClassifier
dt_model=DecisionTreeClassifier(max_depth=5)
dt_model.fit(x_train,y_train)
dt_pred=dt_model.predict(x_test)
dt_pred
```

```
[58]: array([0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
          0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1,
          0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1,
          0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
          0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,
          1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1,
          0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0,
          0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0,
          0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0], dtype=int64)
```

```
[59]: accuracy_score(dt_pred,y_test)
```

```
[59]: 0.7291666666666666
```

```
[60]: confusion_matrix(dt_pred,y_test)
```

```
[60]: array([[100, 22],
           [ 30, 40]], dtype=int64)
```

```
[61]: print("Sensitivity : ",100/(100+40))
```

Sensitivity : 0.7142857142857143

```
[62]: print("Specificity : ",22+(22+30))
```


Specificity : 74

```
[63]: print(classification_report(dt_pred,y_test))
```

	precision	recall	f1-score	support
0	0.77	0.82	0.79	122
1	0.65	0.57	0.61	70
accuracy			0.73	192
macro avg	0.71	0.70	0.70	192
weighted avg	0.72	0.73	0.73	192

```
[64]: dt_pred_proba=dt_model.predict_proba(x_test)[: ,1]
fpr,tpr,thresholds=roc_curve(y_test,svc_pred_proba)
roc_auc=auc(fpr,tpr)
display=RocCurveDisplay(fpr=fpr,tpr=tpr,roc_auc=roc_auc,estimator_name="example_
→estimator")
display.plot()
plt.plot(fpr,fpr,"r--",color="pink")
```

```
[64]: [<matplotlib.lines.Line2D at 0x1ba8fcb4850>]
```



1.12 Random forest

```
[65]: from sklearn.ensemble import RandomForestClassifier
      rf_model=RandomForestClassifier(n_estimators=25)
      rf_model.fit(x_train,y_train)
      rf_pred=rf_model.predict(x_test)
      rf_pred
```

```
[65]: array([0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0,
            0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1,
            0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1,
            1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,
            0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0,
            0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
            0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1], dtype=int64)
```

```
[66]: accuracy_score(rf_pred,y_test)
```

```
[66]: 0.7708333333333334
```

```
[67]: confusion_matrix(rf_pred,y_test)
```

```
[67]: array([[109, 23],
           [ 21, 39]], dtype=int64)
```

```
[68]: print("Sensitivity :",113/(113+41))
```

Sensitivity : 0.7337662337662337

```
[69]: print("Specificity :",21/(21+17))
```


Specificity : 0.5526315789473685

```
[70]: print(classification_report(rf_pred,y_test))
```

	precision	recall	f1-score	support
0	0.84	0.83	0.83	132
1	0.63	0.65	0.64	60
accuracy			0.77	192
macro avg	0.73	0.74	0.74	192
weighted avg	0.77	0.77	0.77	192

```
[71]: rf_pred_proba=rf_model.predict_proba(x_test)[:,-1]
fpr,tpr,thresholds=roc_curve(y_test,svc_pred_proba)
roc_auc=auc(fpr,tpr)
display=RocCurveDisplay(fpr=fpr,tpr=tpr,roc_auc=roc_auc,estimator_name="example_
→estimator")
display.plot()
plt.plot(fpr,fpr,"r--",color="yellow")
```

```
[71]: [<matplotlib.lines.Line2D at 0x1ba8fd8d8b0>]
```



1.13 K nearest neighbors

```
[72]: from sklearn.neighbors import KNeighborsClassifier
      knn_model=KNeighborsClassifier(n_neighbors=5)
      knn_model.fit(x_train,y_train)
      knn_pred=knn_model.predict(x_test)
      knn_pred
```

[illegible]

```
[73]: accuracy_score(knn_pred,y_test)
```

```
[73]: 0.7447916666666666
```

```
[74]: confusion_matrix(knn_pred,y_test)
```

```
[74]: array([[110, 29],
           [ 20, 33]], dtype=int64)
```

```
[75]: print("Sensitivity: ",110/(110+33))
```

Sensitivity: 0.7692307692307693

```
[76]: print("Specificity :",29/(29+20))
```


Specificity : 0.5918367346938775

```
[77]: print(classification_report(knn_pred,y_test))
```

	precision	recall	f1-score	support
0	0.85	0.79	0.82	139
1	0.53	0.62	0.57	53
accuracy			0.74	192
macro avg	0.69	0.71	0.70	192
weighted avg	0.76	0.74	0.75	192

```
[78]: knn_pred_proba=knn_model.predict_proba(x_test)[:,-1]
fpr,tpr,thresholds=roc_curve(y_test,svc_pred_proba)
roc_auc=auc(fpr,tpr)
display=RocCurveDisplay(fpr=fpr,tpr=tpr,roc_auc=roc_auc,estimator_name="example_
→estimator")
display.plot()
plt.plot(fpr,fpr,"r--",color="purple")
```

```
[78]: [<matplotlib.lines.Line2D at 0x1ba8fdf8820>]
```

1.13.1 From the above models, it is clear that Logistic regression and Random forest are best models for this dataset.

1.13.2 THANK YOU...!!!